

Derby Tools and Utilities Guide

Version 10

Derby Document build:
October 2, 2006, 8:04:32 AM (PDT)

Contents

Copyright.....	4
About this guide.....	5
Purpose of this document.....	5
Audience.....	5
How this guide is organized.....	5
What are the Derby tools and utilities?.....	6
Overview.....	6
Environment setup and the Derby tools.....	6
About Derby databases.....	7
JDBC connection basics.....	7
JDBC drivers overview.....	7
Database connection URLs.....	7
Tools and localization.....	8
About locales.....	8
Database territory.....	9
Specifying an alternate codeset.....	9
Formatting display of locale-sensitive data.....	9
Using ij.....	10
Starting ij.....	10
Starting ij using properties.....	10
Getting started with ij.....	11
Connecting to a Derby database.....	11
Using ij commands.....	13
Running ij scripts.....	13
ij properties reference.....	15
ij.connection.connectionName.....	15
ij.database.....	15
ij.dataSource.....	16
ij.driver.....	16
ij.maximumDisplayWidth.....	17
ij.outfile.....	17
ij.password.....	17
ij.protocol.....	17
ij.protocol.protocolName.....	18
ij.showErrorCode.....	18
ij.URLCheck.....	18
ij.user.....	19
derby.ui.codeset.....	20
ij commands and errors reference.....	22
ij commands.....	22
Conventions for ij examples.....	22
ij SQL command behavior.....	22
Absolute.....	23
After Last	23
Async.....	24
Autocommit.....	24
Before First.....	25
Close.....	25
Commit.....	25
Connect.....	26
Describe.....	26
Disconnect.....	27
Driver.....	27
Elapsedtime.....	28
Execute.....	28
Exit.....	29

First	29
Get Cursor	29
Get Scroll Insensitive Cursor	30
Help	31
Last	32
LocalizedDisplay	32
MaximumDisplayWidth	32
Next	33
Prepare	33
Previous	34
Protocol	34
ReadOnly	34
Relative	35
Remove	35
Rollback	36
Run	36
Set Connection	37
Show	37
Wait For	40
Comment	40
Identifier	40
String	41
ij errors	41
ERROR SQLState.....	41
WARNING SQLState.....	41
IJ ERROR.....	42
IJ WARNING.....	42
JAVA ERROR.....	42
Using the bulk import and export procedures	43
Bulk Import/Export overview	43
Options for running the import and export procedures.....	43
Bulk import/export requirements and considerations.....	43
Bulk import and export	44
Bulk-Import.....	44
Bulk-Export.....	45
Examples of bulk import and export.....	46
Importing into tables with identity columns.....	47
Executing import/export procedures from JDBC.....	48
File format for input and output.....	48
Treatment of NULLS.....	49
CODESET values for import/export.....	49
Storing jar files in a database	51
Adding a Jar File	51
Removing a jar file	51
Replacing a jar file	51
Installing a jar example	52
sysinfo	53
sysinfo example	53
Using sysinfo to check the classpath	53
dblook	55
Using dblook	55
dblook options	55
Generating the DDL for a database	56
dblook examples	56
Trademarks	58

Copyright

Copyright 2004, 2006 The Apache Software Foundation or its licensors, as applicable.

Licensed to the Apache Software Foundation (ASF) under one or more contributor license agreements. See the NOTICE file distributed with this work for additional information regarding copyright ownership. The ASF licenses this file to you under the Apache License, Version 2.0 (the "License"); you may not use this file except in compliance with the License. You may obtain a copy of the License at

<http://www.apache.org/licenses/LICENSE-2.0>

Unless required by applicable law or agreed to in writing, software distributed under the License is distributed on an "AS IS" BASIS, WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied. See the License for the specific language governing permissions and limitations under the License.

About this guide

For general information about the Derby documentation, such as a complete list of books, conventions, and further reading, see *Getting Started with Derby*.

Purpose of this document

This book describes how to use the Derby tools and utilities. The tools and utilities covered in this book include:

- ij
- the import and export utilities
- the database class loading utilities
- sysinfo
- dblook

Audience

This book is for:

- developers, who might use the tools when developing applications
- system administrators, who might use the tools to run backup scripts or to import large amounts of data
- end-users, who might use one of the tools to run ad-hoc queries against a database

How this guide is organized

This guide includes the following sections:

- [What are the Derby tools and utilities?](#)

Overview of the tools and utilities, and Derby and JDBC basics for new or infrequent users.

- [Using ij](#)

How to get started with ij, a JDBC and SQL scripting tool.

- [ij properties reference](#)

Reference for ij properties.

- [ij commands and errors reference](#)

Reference for ij commands and errors.

- [Using the bulk import and export procedures](#)

Reference and how-to instructions for using bulk import and export.

- [ij commands and errors reference](#)

Syntax for executing the built-in procedures for storing jar files in the database.

- [sysinfo](#)

Reference information on the utility that provides information about your Derby environment.

- [dblook](#)

dblook is Derby's Data Definition Language (DDL) Generation Utility, which is more informally called a schema dump tool. It is a simple utility that dumps the DDL of a user-specified database to either a console or to a file. The generated DDL can then be used to recreate all or parts of a database, view a subset of a database's objects (for example, those which pertain to specific tables or schemas), or document a database's schema.

What are the Derby tools and utilities?

The Derby tools and utilities are a set of routines supplied with Derby that are typically used to setup and update a Derby database.

For more complete information on developing a system using Derby, see the *Derby Developer's Guide*.

Overview

Derby is a database management system (DBMS), accessed by applications through the JDBC API.

Included with the product are some standalone Java tools and utilities that make it easier to use and develop applications for Derby.

These tools and utilities include:

- *ij*

ij is Derby's interactive JDBC scripting tool. It is a simple utility for running scripts against a Derby database. You can also use it interactively to run ad hoc queries. *ij* provides several commands for ease in accessing a variety of JDBC features.

ij can be used in an embedded or a client/server environment.

- *The import and export utilities*

These server-side utilities allow you to import data directly from files into tables and to export data from tables into files. Server-side utilities can be in a client/server environment but require that all files referenced be on the Server machine.

- *Database class loading utilities*

These utilities allow you to store application logic in a database.

- *sysinfo*

sysinfo provides information about your version of Derby and your environment.

- *dblook*

dblook is Derby's Data Definition Language (DDL) Generation Utility, also called a schema dump tool. It is a simple utility for the dumping the DDL of a user-specified database to either a console or to a file. The generated DDL can then be used for such things as recreating all or parts of a database, viewing a subset of a database's objects (for example, those which pertain to specific tables and schemas), or documenting a database's schema.

Environment setup and the Derby tools

ij, *sysinfo*, and *dblook* are tools that can be used in an embedded or a client/server environment. The import and export utilities and database class loading utilities are database-side utilities, which means that they run in the same JVM as Derby (e.g. on the Server). This means when used in a client/server environment all files imported, exported, or loaded must be local to the Server machine.

Java 2 Platform, Standard Edition, Version 1.3

All Derby tools require the Java 2 Platform, Standard Edition, Version 1.3 or later.

Classpath

To simplify the process of setting up the *CLASSPATH* environment variable to run Derby and the tools a new jarfile, *derbyrun.jar*, has been added to the Derby distribution. This is the only jarfile that needs to be in your classpath. The following information is provided for people using Derby distributions that do not contain *derbyrun.jar* or who prefer to setup the classpath manually.

The following lists the files that need to be in the Derby classpath for each tool and configuration:

- To use *ij*, you must have *derbytools.jar* in your classpath. If you are using the embedded driver, you must also include *derby.jar*.
- To use the import and export utilities and the database class loading utilities, you must have *derby.jar* in your classpath.
- To use *sysinfo*, either *derby.jar* or *derbytools.jar* must be in your classpath.
- To use Derby tools from a client with the Derby Network Server, you must have *derbyclient.jar* and *derbytools.jar* in your classpath. See the *Derby Server and Administration Guide* for more information.

About Derby databases

A Derby database consists of platform-independent files stored in a directory that has the same name as the database.

JDBC connection basics

Most of the Derby tools are JDBC applications. A JDBC application is one that uses the classes in the *java.sql* package to interact with a DBMS.

When you work with JDBC applications, you need to know about several concepts. The most basic is the *connection*. A *JDBC connection* is the object through which commands are sent to the Derby engine and responses are returned to the program. Establishing a *connection* to a specific database is done by specifying a appropriate database *URL*. The following sections provide background information to help in understanding the Derby database connection URL.

JDBC drivers overview

Before a JDBC application connects to a database, it must cause the proper JDBC driver to be loaded in the Java session. Derby provides the following JDBC drivers for use with the Derby database engine:

- *org.apache.derby.jdbc.EmbeddedDriver*

For embedded environments, when Derby runs in the same JVM as the application. This is commonly referred to as the embedded driver.

- *org.apache.derby.jdbc.ClientDriver*

For client/server environments that use the Derby Network Server. This is commonly referred to as the Network Client driver.

You can use *ij* to connect to any database that supplies a JDBC driver. For those databases, you would need to load the supplied JDBC driver.

Database connection URLs

A JDBC URL provides a way of identifying a database so that the appropriate driver recognizes it and connects to it. In the Derby documents, a JDBC URL is referred to as a

database connection URL.

After the driver is loaded, an application must specify the correct database connection URL to connect to a specific database. The Derby database connection URL allows you to accomplish tasks other than simply connecting. For more information about the Derby database connection URLs, see the *Derby Developer's Guide*.

A JDBC URL always starts with *jdbc:*. After that, the format for the database connection URL depends on the JDBC driver.

Here is the format for the database connection URL for connecting to an existing Derby database using the embedded driver:

- *jdbc:derby:databaseName;URLAttributes*

The format for the database connection URL for connecting to an existing Derby database using the Network Client is:

- *jdbc:derby://host:port/databaseName;URLAttributes*

The italicized items stand for something the user fills in:

- *databaseName*

The name of the database you want to connect to. This might also include the file system path to the database.

- *URLAttributes*

One or more of the supported attributes of the database connection URL, such as *;upgrade=true*, *create=true* or *territory=ll_CC*. For more information, see the *Derby Developer's Guide*.

- *host*

The name of the machine where the server is running. It can be the name of the machine or the address.

- *port*

The port number used by the server framework

About Protocols

Officially, the portion of the database connection URL called the protocol is *jdbc:*, just as *http://* is a protocol in Web URLs. However, the second portion of the database connection URL (everything between *jdbc:* and *databaseName*), which is called the subprotocol, is informally considered part of the protocol. Later in this book you might see references to protocol. Consider protocol to be everything that comes before *databaseName*.

For complete information about the database connection URL, see the *Derby Developer's Guide*.

Tools and localization

The Derby tools provide support for common localization features such as localized message files and GUI, locale-appropriate formatting of data, codesets, unicode identifiers and data, and database territories.

For general information about international Derby systems, see the *Derby Developer's Guide*.

About locales

In the Derby documentation, we refer to three locales:

- *Java System locale*

This is the locale of your machine, which is automatically detected by your JVM. For Derby and Derby tools, the Java system locale determines the default locale.

- *Database territory*

This is the territory associated with your database when it is created. By default, this is the same as the [java system locale](#). The database territory determines the language of database errors.

- *ij or dblook Session locale*

This locale is associated with your `ij` or `dblook` session. This locale determines the localized display format for numbers, dates, times, and timestamps.

Database territory

To specify a database territory, use the *territory* attribute on the URL connection when creating the database.

Note: You cannot modify a database's territory after the database has been created.

For information about database territories, see the Internationalization appendix in the *Derby Developer's Guide*.

Specifying an alternate codeset

You can specify an alternate codeset for your tool session.

Use the [derby.ui.codeset](#) property when starting `ij` or `dblook`. This property can be useful when working with scripts created on a different system.

Formatting display of locale-sensitive data

To display dates, timestamps, numbers, and times in the format of the `ij` Session locale, use the [LocalizedDisplay](#) command.

Note: These options do not change how Derby stores locale-sensitive data, simply how the tool displays the data.

The following example demonstrates using *localizedDisplay* in an *en_US* locale:

```
ij> VALUES CURRENT_DATE;
1
-----
2001-09-06
1 row selected
ij> localizeddisplay on;
ij> VALUES CURRENT_DATE;
1
-----
September 6, 2001
1 row selected
```

Using ij

ij is Derby's interactive JDBC scripting tool. It is a simple utility for running scripts or interactive queries against a Derby database.

ij is a Java application, which you start from a command window such as an MS-DOS Command Window or the UNIX shell. ij provides several non-SQL commands for ease in accessing a variety of JDBC features for testing.

Starting ij

Derby provides batch and shell scripts for users in Windows and UNIX environments that can be used to start IJ. By calling the appropriate script you will start ij and be able to connect with a simple command. These scripts set the [ij.protocol](#) property, which simplifies the process of connecting to a database by eliminating the need to specify the protocol portion of the connection URL. The scripts are found in the %DERBY_HOME%/frameworks/ directory tree. You can also customize the ij scripts to suit your environment.

If you are starting ij from a command line, be sure that derbyrun.jar is in your classpath (for pre-10.2 distributions derbytools.jar and usually derby.jar were required in the classpath). If you are using Derby as a client/server environment, start the server before connecting to the Derby database. You can start ij by running the ij scripts for your environment in the %DERBY_HOME%/frameworks/embedded/bin/ directory or in the %DERBY_HOME%/frameworks/NetworkServer/bin/ directory.

To start ij, run the script provided or use this command:

```
java [<options>] org.apache.derby.tools.ij
    [-p <propertyFile>] [<inputFile>]
```

The command line items are:

- java

Start the JVM.

- options

The options that the JVM uses. You can use the -D command to set ij properties (see [Starting ij using properties](#)) or system properties, such as Derby properties.

- propertyFile

A file you can use to set ij properties (instead of the -D command). The property file should be in the format created by the java.tools.Properties.save methods, which is the same format as the derby.properties file.

- inputFile

A file from which to read commands. The ij tool exits at the end of the file or an exit command. Using an input file causes ij to print out the commands as it runs them. If you reroute standard input, ij does not print out the commands. If you do not supply an input file, ij reads from the standard input.

For detailed information about ij commands, see [ij commands and errors reference](#).

Starting ij using properties

You set ij properties in any of the following ways:

1. by using the `-D` command on the command line
2. by specifying a properties file using the `-p propertyfile` option on the command line

Remember: ij property names are case-sensitive, while commands are case-insensitive.

The following examples illustrate how to use ij properties:

To start ij by using a properties file called `ij.properties`, use the following command:

```
java org.apache.derby.tools.ij -p ij.properties
```

To start ij with a `maximumDisplayWidth` of 1000:

```
java -Dij.maximumDisplayWidth=1000 org.apache.derby.tools.ij
```

To start ij with an `ij.protocol` of `jdbc:derby:` and an `ij.database` of `sample`, use the following command:

```
java -Dij.protocol=jdbc:derby: -Dij.database=sample
org.apache.derby.tools.ij
```

To start ij with two named connections, using the `ij.connection.connectionName` property, use the following command:

```
java -Dij.connection.sample=jdbc:derby:sample
-Dij.connection.History=jdbc:derby:History
-Dderby.system.home=c:\derby\demo\databases
org.apache.derby.tools.ij
```

To see a list of connection names and the URL's used to connect to them, use the following command: (If there is a connection that is currently active, it will show up with an `*` after its name.)

```
ij version 10.2
ij(HISTORY)> show connections;
HISTORY* -      jdbc:derby:History
SAMPLE -      jdbc:derby:sample
* = current connection
ij(HISTORY)>
```

Getting started with ij

This section discusses the use of the ij tool.

Connecting to a Derby database

To connect to a Derby database, you need to perform the following steps:

1. Start the JVM
2. Load the appropriate driver.
3. Create a connection by providing a valid database connection URL.

When using ij interactively to connect to a Derby database connection information is generally supplied on the full database connection URL. ij automatically loads the appropriate driver based on the syntax of the URL. The following example shows how to connect in this manner by using the **Connect** command and the embedded driver:

```
D:>java org.apache.derby.tools.ij
ij version 10.2
ij> connect 'jdbc:derby:sample';
ij>
```

If the URL entered contains Network Client information the **Connect** command loads the Network Client driver:

```
D:>java org.apache.derby.tools.ij
ij version 10.2
ij> connect 'jdbc:derby://localhost:1527/sample';
ij>
```

Note: In these and subsequent examples the databases were created in the `derby.system.home` directory. For more information on the System Directory see the *Derby Developer's Guide*.

ij provides alternate methods of specifying part or all of a connection URL (e.g. the [ij.protocol](#), [ij.database](#), or [ij.connection.connectionName](#) properties). These properties are often used when a script is being used and the path to the database or the driver name is not known until runtime. The properties can also be used to shorten the amount of information that must be provided with the connection URL. The following are some examples of different ways to supply the connection information:

- Supplying full connection information on the command line
Specifying one of the following properties along with a valid connection URL on the ij command line starts ij with the connection already active. This is often used when running a SQL script so the database name or path can be specified at runtime.
 - [ij.database](#) - opens a connection using the URL provided
 - [ij.connection.connectionName](#) - Used to open one or more connections. The property can appear multiple times on the command line with different *connectionNames* and the same or different URLs.

This example shows how to create the database *myTours* and run the script *ToursDB_schema.sql* by specifying the database URL using the [ij.database](#) property.

```
C:\>java -Dij.database=jdbc:derby:myTours;create=true
org.apache.derby.tools.ij
%DERBY_HOME%\demo\programs\toursdb\ToursDB_schema.sql
ij version 10.2
CONNECTION0* - jdbc:derby:myTours
* = current connection
ij> -- Licensed to the Apache Software Foundation (ASF) under one or
more
-- contributor license agreements. See the NOTICE file distributed
with
...output removed...
ij> CREATE TRIGGER TRIG2 AFTER DELETE ON FLIGHTS FOR EACH STATEMENT
MODE DB2SQL
INSERT INTO FLIGHTS_HISTORY (STATUS) VALUES ('INSERTED FROM TRIG2');
0 rows inserted/updated/deleted
ij>
```

- Defining a Protocol and using a "short form" URL

A default URL protocol and subprotocol can be specified by setting the property [ij.protocol](#) or using the [ij.Protocol](#) command. This allows a connection to be made by specifying only the database name. This "short form" of the database connection URL defaults the protocol (For more information, see [About Protocols](#)).

This example uses the [ij.Protocol](#) command and a "short form" connection URL:

```
D:>java org.apache.derby.tools.ij
ij version 10.2
ij> protocol 'jdbc:derby:';
ij> connect 'sample';
ij>
```

- Specifying an alternate Driver

If you are using the drivers supplied by Derby, you can specify the driver names listed in [JDBC drivers overview](#). However, the Derby drivers are implicitly loaded when a supported protocol is used so specifying them is probably redundant. Specifying a driver is required when ij is used with other JDBC drivers to connect to non-Derby databases. To use drivers supplied by other vendors explicitly specify the driver one of three ways

- with an ij property [ij.Driver](#)
- using the JVM system property `jdbc.drivers`
- using the ij [Driver](#) command

This example specifies the driver using the ij [Driver](#) command

```
D:>java org.apache.derby.tools.ij
ij>version 10.2
ij>driver 'sun.jdbc.odbc.JdbcOdbcDriver';
ij>connect 'jdbc:odbc:myOdbcDataSource';
ij>
```

The ij Driver name and connection URL

[Specifying the Driver Name and database connection URL](#), summarizes the different ways to specify the driver name and database connection URL.

Table 1. Specifying the Driver Name and database connection URL

Action	System Property	ij Property	ij Command
loading the driver implicitly		ij.connection.connectionName (plus full URL) ij.database (plus full URL) ij.protocolij.protocol.protocolName (plus protocol clause in Connect command)	ProtocolConnect (plus full URL)
loading the driver explicitly	<code>jdbc.drivers</code>	<code>-Dij.Driver</code>	Driver
specifying the database connection URL		ij.connection.connectionName	Connect

Using ij commands

The primary purpose of ij is to allow the execution of Derby SQL statements interactively or via scripts. Since SQL statements can be quite long ij uses the semicolon to mark the end of statement or command. All statements must be terminated with a semicolon. ij uses properties, listed later in this guide, to simplify its use. ij also recognizes specialized commands that extends its functionality such as the ability to create and test cursors and prepared statements, transaction control and more. For complete information about ij commands, see [ij commands and errors reference](#) later in the Tools Guide.

Other uses for ij

ij is a JDBC-neutral scripting tool with a small command set. It can be used to access any JDBC driver and database accessible through that driver.

The main benefit of a tool such as ij is that it is easy to run scripts for creating a database schema and automating other repetitive database tasks.

In addition, ij accepts and processes SQL commands interactively for ad hoc database access.

Running ij scripts

You can run scripts in ij in any of the following ways:

- Name an input file as a command-line argument.

For example:

```
java org.apache.derby.tools.ij <myscript.sql>
```

- Redirect standard input to come from a file.

For example:

```
java org.apache.derby.tools.ij <myscript.sql>
```

- Use the [Run](#) command from the ij command line.

For example:

```
ij> run 'myscript.sql';
```

Note: If you name an input file as a command-line argument or if you use the [Run](#) command, ij echoes input from a file. If you redirect standard input to come from a file, ij does not echo commands.

You can save output in any of the following ways:

- By redirecting output to a file:

```
java org.apache.derby.tools.ij <myscript.sql> > <myoutput.txt>
```

- By setting the `ij.outfile` property:

```
java -Dij.outfile=<myoutput.txt> org.apache.derby.tools.ij  
    <myscript.sql>
```

ij exits when you enter the [Exit](#) command or, if executing a script, when the end of the command file is reached. When you use the [Exit](#) command, ij automatically shuts down an embedded Derby system by issuing a `connect jdbc:derby:;shutdown=true` request. It does not shut down Derby if it is running in a server framework.

ij properties reference

When starting up `ij`, you can specify properties on the command line or in a properties file, as described in [Starting ij using properties](#).

ij.connection.connectionName

Function

Creates a named connection to the given database connection URL when `ij` starts up; it is equivalent to the Connect AS *Identifier* command. The database connection URL can be of the short form if an [ij.protocol](#) is specified. This property can be specified more than once per session, creating multiple connections. When `ij` starts, it displays the names of all the connections created in this way. It also displays the name of the current connection, if there is more than one, in the `ij` prompt.

Syntax

```
ij.connection.connectionName=databaseConnectionURL
```

When specified on the command line the *databaseConnectionURL* should not be enclosed in single quotations, however, if the database path contains special characters (e.g. a space) it must be enclosed in double quotes.

Example

This example connects to the existing database *sample* and creates then connects to the database *anotherDB*.

```
D:> java -Dij.connection.sample1=jdbc:derby:sample
-Dij.connection.anotherConn=jdbc:derby:anotherDB;create=true
org.apache.derby.tools.ij
ij version 10.2
ANOTHERCONN* - jdbc:derby:anotherDB;create=true
SAMPLE1 - jdbc:derby:sample
* = current connection
ij(ANOTHERCONN)>
```

See also

- [Connect](#)

ij.database

Function

Creates a connection to the database name listed indicated by the property when `ij` starts up. You can specify the complete connection URL (including protocol) with this property or just the database name if you also specify *ij.protocol* on the command line. After it boots, `ij` displays the generated name of the connection made with this property.

Syntax

```
ij.database=databaseConnectionURL
```

When specified on the command line the *databaseConnectionURL* should not be enclosed in single quotations, however, if the database path contains special characters (e.g. a space) it must be enclosed in double quotes.

Example

```
java -Dij.protocol=jdbc:derby:
-Dij.database=wombat;create=true org.apache.derby.tools.ij
```

```
ij version 10.2
CONNECTION0* - jdbc:derby:wombat
* = current connection
ij>
```

ij.dataSource

Attention: DerbyVersion 10.2 includes JDBC 4.0 functionality based on a pre-release of Java SE 6. When you use Derby with a Java SE 6 virtual machine, Derby will use its JDBC 3.0 drivers, just as with 1.4 and 1.5 virtual machines. JDBC 4.0 is available only to developers who download Java SE 6 and use it to build in support for JDBC 4.0. The Release Notes explain how to build the JDBC 4.0 functionality with DerbyVersion 10.2. A follow-on Derby release will include JDBC 4.0 functionality based on the final JDBC 4.0 specification once Java SE 6 becomes generally available.

Function

The `ij.dataSource` property specifies the datasource to be used to access the database. When specifying a datasource, Derby does not use the DriverManager mechanism to establish connections.

Syntax

To establish a connection using `ij.dataSource`, you need to set the `ij.dataSource.databaseName` and to create the database specify the `ij.dataSource.createDatabase` property as well.

```
ij.dataSource=datasourcenameij.dataSource.databaseName=databasename[ij.dataSource.createDatabase=createDatabaseName]
```

Example

```
#
# If your application runs on JDK 1.6 or higher, then you should
# specify the JDBC4 variant of this DataSource:
# org.apache.derby.jdbc.EmbeddedDataSource40.
#
java -Dij.dataSource=org.apache.derby.jdbc.EmbeddedDataSource
-Dij.dataSource.databaseName=sample -Dij.dataSource.createDatabase=create
org.apache.derby.tools.ij
ij version 10.2
CONNECTION0*
* = current connection
```

For more information about DataSources, refer to the JDBC documentation and "Using Derby as a J2EE Resource Manager" in the *Derby Developer's Guide*.

ij.driver

Function

Loads the JDBC driver that the class specifies.

Syntax

```
ij.driver=JDBCDriverClassName
```

Notes

Example

```
D:>java -Dij.driver=sun.jdbc.odbc.JdbcOdbcDriver
org.apache.derby.tools.ij
ij version 10.2
ij> Connect 'jdbc:odbc:MyODBCDataSource';
ij>
```

See also

- [Driver](#)

ij.maximumDisplayWidth

Function

Specifies the maximum number of characters used to display any column. The default value is 128. Values with display widths longer than the maximum are truncated and terminated with an & character.

Syntax

```
ij.maximumDisplayWidth=numberOfCharacters
```

Example

```
java -Dij.maximumDisplayWidth=1000 org.apache.derby.tools.ij
```

See also

- [MaximumDisplayWidth](#)

ij.outfile

Function

Specifies a file to which the system should direct output for a session. Specify the file name relative to the current directory, or specify the absolute path.

Syntax

```
ij.outfile=fileName
```

Example

```
java -Dij.outfile=out.txt org.apache.derby.tools.ij myscript.sql
```

ij.password

Function

Specifies the password used to make connections. This property is used in conjunction with the *ij.user* property to authenticate a connection. If authentication is not active then these properties are ignored.

Syntax

```
ij.password=password
```

Example

```
java -Dij.user=me -Dij.password=mime org.apache.derby.tools.ij
```

See the *Derby Developer's Guide* for more information on Derby authentication and security.

ij.protocol

Function

Specifies the default protocol and subprotocol portions of the database connection URL for connections. The Derby protocol is:

- jdbc:derby:

Allows you to use a short form of a database name in a connection URL.

Syntax

```
ij.protocol=protocolForEnvironment
```

Example

```
D:>java -Dij.protocol=jdbc:derby:
      org.apache.derby.tools.ij
ij version 10.2
ij> Connect 'newDB;create=true';
ij>
```

See also

- [Protocol](#)

ij.protocol.protocolName

Function

This property is similar to the [ij.protocol](#) property. The only difference is that it associates a name with the value, thus allowing you to define and use more than one protocol. (See [Connect](#).)

Syntax

```
ij.protocol.protocolName=protocolForEnvironment
```

Example

```
D:>java -Dij.protocol.derby=jdbc:derby:
-Dij.protocol.emp=jdbc:derby: org.apache.derby.tools.ij
ij version 10.2
ij> Connect 'newDB' protocol derby as new;
ij>
```

See also

- [Protocol](#)

ij.showErrorCode

Function

Set this property to *true* to have ij display the *SQLException ErrorCode* value with error messages. The default is *false*.

Error codes denote the severity of the error. For more information, see the *Derby Reference Manual*.

Syntax

```
ij.showErrorCode=trueOrFalse
```

Example

```
java -Dij.showErrorCode=true -Dij.protocol=jdbc:derby:
      org.apache.derby.tools.ij
ij version 10.2
ij> Connect 'sample';
ij> VLUES 1;
ERROR 42X01: Syntax error: Encountered "VLUES"
at line 1, column 1. (errorCode = 30000)
ij>
```

ij.URLCheck

Function

This property determines whether `ij` checks for invalid or non-Derby URL attributes. Set this property to *false* to prevent `ij` from validating URL attributes. The default value is *true*.

When the *ij.URLCheck* property is set to *true*, you are notified whenever a connection URL contains an incorrectly specified attribute. For example if the attribute name is misspelled or cased incorrectly `ij` prints a message.

Note: `ij` checks attribute *values* if the attribute has pre-defined values. For example, the attribute *unicode* has the pre-defined values of *true* or *false*. If you try to set the attribute *unicode* to a value other than *true* or *false*, `ij` displays an error. For example:

```
ij> Connect 'jdbc:derby:anyDB;create=true;unicode=falj';
ERROR XJ05B: JDBC attribute 'unicode' has an invalid value 'falj',
valid values are '{true|false}'.
ij>
```

Syntax

```
ij.URLCheck={ false | true }
```

Example

By default, `ij` displays messages about invalid attributes:

```
java org.apache.derby.tools.ij
ij version 10.2
ij> connect 'mydb;uSer=naomi';
URL Attribute [uSer=naomi]
Case of the Derby attribute is incorrect.
```

The following command line specifies to turn off URL attribute checking in `ij`.

```
java -Dij.URLCheck=false org.apache.derby.tools.ij
ij version 10.2
ij> connect 'mydb;uSer=naomi';
ij>
```

Typically, you would only explicitly turn off the URL checker if you were using `ij` with a non-Derby JDBC driver or database.

Notes

URLCheck does not check the validity of properties, only database connection URL *attributes*.

`ij` recognizes the following attributes:

- *bootPassword*
- *create*
- *databaseName*
- *dataEncryption*
- *encryptionAlgorithm*
- *encryptionProvider*
- *territory*
- *logDevice*
- *password*
- *shutdown*
- *unicode*
- *upgrade*
- *user*

ij.user

Function

Specifies the logon name used to establish the connection. This property is used in conjunction with the *ij.password* property to authenticate a connection. If authentication is not active then these properties are ignored.

When a username is supplied you need to be aware of the database schema. When you connect using *ij.user*, the default database schema applied to all SQL statements is the same as the user id provided even if the schema does not exist. Use the SET SCHEMA statement to change the default when the schema does not match the username. Alternately you can fully qualify the database objects referred to in the SQL statements. If no user is specified, no SET SCHEMA statement has been issued, or SQL statements do not include the schema name, all database objects are assumed to be under the APP schema.

Syntax

```
ij.user=username
```

Example

```
java -Dij.user=me -Dij.password=mime org.apache.derby.tools.ij
ij version 10.2
ij> connect 'jdbc:derby:sampleDB';
ij> set schema finance;
ij> select * from accounts;
```

See the *Derby Developer's Guide* for more information on Derby and security.

derby.ui.codeset

Function

Set this property to a supported character encoding value when using one of the Derby tools with a language not supported by your default system.

Syntax

```
derby.ui.codeset=derbyval
```

where *derbyval* is a supported character encoding value, for example, UTF8 (see [Sample Character Encodings](#)).

Example

The following command line specifies to run *ij* using the Japanese territory (*territory=ja_JP*) using Japanese Latin Kanji mixed encoding (*codeset=Cp939*):

```
java -Dderby.ui.territory=ja_JP -Dderby.ui.codeset=Cp939
-Dij.protocol=jdbc:derby:
org.apache.derby.tools.ij
```

The following table contains a sampling of character encodings supported by the IBM Application Developer Kit. To see the full list, go to

<http://java.sun.com/j2se/1.3/docs/guide/intl/encoding.doc.html>.

Table 2. Sample Character Encodings

Character Encoding	Explanation
8859_1	ISO Latin-1

Character Encoding	Explanation
8859_2	ISO Latin-2
8859_7	ISO Latin/Greek
Cp1257	Windows Baltic
Cp1258	Windows Vietnamese
Cp437	PC Original
EUCJIS	Japanese EUC
GB2312	GB2312-80 Simplified Chinese
JIS	JIS
KSC5601	KSC5601 Korean
MacCroatian	Macintosh Croatian
MacCyrillic	Macintosh Cyrillic
SJIS	PC and Windows Japanese
UTF8	Standard UTF-8

ij commands and errors reference

This section describes the commands and errors within the ij tool.

ij commands

ij accepts several commands to control its use of JDBC. It recognizes a semicolon as the end of an ij or SQL command; it treats semicolons within SQL comments, strings, and delimited identifiers as part of those constructs, not as the end of the command. Semicolons are required at the end of an ij or SQL statement.

All ij commands, identifiers, and keywords are case-insensitive.

Commands can span multiple lines without any special escaping for the ends of lines. This means that if a string spans a line, the new lines will show up in the value in the string.

ij treats any command that it does not recognize as an SQL command to be passed to the underlying connection, so syntactic errors in ij commands will cause them to be handed to the SQL engine and will probably result in SQL parsing errors.

Conventions for ij examples

Examples in this document show input from the keyboard or a file in bold text and console output from the DOS prompt or the ij application in regular text.

```
C:\> REM This example is from a DOS prompt:
C:\> java -Dij.protocol=jdbc:derby: org.apache.derby.tools.ij
ij version 10.2
ij> connect 'menuDB;create=true';
ij> CREATE TABLE menu(course CHAR(10), item CHAR(20), price INTEGER);
0 rows inserted/updated/deleted
ij> disconnect;
ij> exit;
C:\>
```

ij SQL command behavior

Any command other than those documented in the ij command reference are handed to the current connection to execute directly. The statement's closing semicolon, used by ij to determine that it has ended, is not passed to the underlying connection. Only one statement at a time is passed to the connection. If the underlying connection itself accepts semicolon-separated statements (which Derby does not), they can be passed to the connection using ij's Execute command to pass in a command string containing semicolon-separated commands.

ij uses the result of the JDBC execute request to determine whether it should print a number-of-rows message or display a result set.

If a JDBC execute request causes an exception, it displays the *SQLState*, if any, and error message.

Setting the ij property [ij.showErrorCode](#) to *true* displays the *SQLException*'s error code (see [ij properties reference](#)).

The number-of-rows message for inserts, updates, and deletes conforms to the JDBC specification for any SQL statement that does not have a result set. DDL (data definition language) commands typically report "0 rows inserted/updated/deleted" when they successfully complete.

To display a result set, `ij` formats a banner based on the JDBC *ResultSetMetaData* information returned from *getColumnLabel* and *getColumnWidth*. Long columns wrap the screen width, using multiple lines. An `&` character denotes truncation (`ij` limits displayed width of a column to 128 characters by default; see [MaximumDisplayWidth](#)).

`ij` displays rows as it fetches them. If the underlying DBMS materializes rows only as they are requested, `ij` displays a partial result followed by an error message if there is a error in fetching a row partway through the result set.

`ij` verifies that a connection exists before issuing statements against it and does not execute SQL when no connection has yet been made.

There is no support in `ij` for the JDBC feature multiple result sets.

ij command example

```
ij> INSERT INTO menu VALUES ('appetizer','baby greens',7),
('entree','lamb chops ',6),('dessert','creme brulee',14);
3 rows inserted/updated/deleted
ij> SELECT * FROM menu;
COURSE      | ITEM                      | PRICE
-----|-----|-----
entree      | lamb chop                 | 14
dessert     | creme brulee              | 6
appetizer   | baby greens               | 7
3 rows selected
ij>
```

Absolute

Syntax

```
ABSOLUTE int Identifier
```

Description

Moves the cursor to the row specified by the *int*, then fetches the row. The cursor must have been created with the [Get Scroll Insensitive Cursor](#) command. It displays a banner and the values of the row.

Note: This command works only in a Java 2 Platform, Standard Edition, v 1.2 (J2SE) or higher environment.

Example

```
ij> autocommit off;
ij> get scroll insensitive cursor scrollCursor as
'SELECT * FROM menu FOR UPDATE OF price';
ij> absolute 3 scrollCursor;
COURSE      | ITEM                      | PRICE
-----|-----|-----
entree      | lamb chop                 | 14
```

After Last

Syntax

```
AFTER LAST Identifier
```

Description

Moves the cursor to after the last row, then fetches the row. (Since there is no current row, it returns the message: `No current row.`")

The cursor must have been created with the [Get Scroll Insensitive Cursor](#) command.

Note: This command works only in a Java 2 Platform, Standard Edition, v 1.2 (J2SE) or

higher environment.

Example

```
ij> get scroll insensitive cursor scrollCursor as
'SELECT * FROM menu FOR UPDATE OF price';
ij> after last scrollcursor;
No current row
```

Async

Syntax

```
ASYNC IdentifierString
```

Description

The ASYNC command lets you execute an SQL statement in a separate thread. It is used in conjunction with the [Wait For](#) command to get the results.

You supply the SQL statement, which is any valid SQL statement, as a *String*. The *Identifier* you must supply for the async SQL statement is used in the [Wait For](#) command and is a case-insensitive `ij` identifier; it must not be the same as any other identifier for an async statement on the current connection. You cannot reference a statement previously prepared and named by the `ijPrepare` command in this command.

`ij` creates a new thread in the current connection to issue the SQL statement. The separate thread is closed once the statement completes.

Example

```
ij> async aInsert 'INSERT into menu values ('entree','chicken',11)';
ij> INSERT INTO menu VALUES ('dessert','ice cream',3);
1 rows inserted/updated/deleted.
ij> wait for aInsert;
1 rows inserted/updated/deleted.
-- the result of the asynchronous insert
```

Autocommit

Syntax

```
AUTOCOMMIT { ON | OFF }
```

Description

Turns the connection's auto-commit mode on or off. JDBC specifies that the default auto-commit mode is `ON`. Certain types of processing require that auto-commit mode be `OFF`. For information about auto-commit, see the *Derby Developer's Guide*.

If auto-commit mode is changed from off to on when there is a transaction outstanding, that work is committed when the current transaction commits, not at the time auto-commit is turned on. Use [Commit](#) or [Rollback](#) before turning on auto-commit when there is a transaction outstanding, so that all prior work is completed before the return to auto-commit mode.

Example

```
ij> autocommit off;
ij> DROP TABLE menu;
0 rows inserted/updated/deleted
ij> CREATE TABLE menu (course CHAR(10), item CHAR(20), price INT);
0 rows inserted/updated/deleted
ij> INSERT INTO menu VALUES ('entree', 'lamb chop', 14),
```

```

('dessert', 'creme brulee', 6),
('appetizer', 'baby greens', 7));
3 rows inserted/updated/deleted
ij> commit;
ij> autocommit on;
ij>

```

Before First

Syntax

```
BEFORE FIRST int Identifier
```

Description

Moves the cursor to before the first row, then fetches the row. (Since there is no current row, it returns the message `No current row.`)

The cursor must have been created with the [Get Scroll Insensitive Cursor](#) command.

Note: This command works only in a Java 2 Platform, Standard Edition, v 1.2 (J2SE) or higher environment.

Example

```

ij> get scroll insensitive cursor scrollCursor as
'SELECT * FROM menu FOR UPDATE OF price';
ij> before first scrollcursor;
No current row

```

Close

Syntax

```
CLOSE Identifier
```

Description

Closes the named cursor. The cursor must have previously been successfully created with the `ij` [Get Cursor](#) or [Get Scroll Insensitive Cursor](#) commands.

Example

```

ij> get cursor menuCursor as 'SELECT * FROM menu';
ij> next menuCursor;
COURSE      |ITEM                |PRICE
-----|-----
entree      |lamb chop           |14
ij> next menuCursor;
COURSE      |ITEM                |PRICE
-----|-----
dessert     |creme brulee        |6
ij> close menuCursor;
ij>

```

Commit

Syntax

```
COMMIT
```

Description

Issues a *java.sql.Connection.commit* request. Use this command only if auto-commit is off. A *java.sql.Connection.commit* request commits the currently active transaction and initiates a new transaction.

Example

```
ij> commit;
ij>
```

Connect

Syntax

```
CONNECT ConnectionURLString [ PROTOCOL Identifier ]
      [ AS Identifier ] [ USER String
      PASSWORD String ]
```

Description

Connects to the database indicated by the *ConnectionURLString*. Specifically, takes the value of the string database connection URL and issues a `java.sql.DriverManager.getConnection` request to set the current connection to that database connection URL.

You have the option of specifying a name for your connection. Use the [Set Connection](#) command to switch between connections. If you do not name a connection, the system generates a name automatically.

You also have the option of specifying a named protocol previously created with the [Protocol](#) command or the `ij.protocol.protocolName` property.

If the connection requires a user name and password, supply those with the optional user and password parameters.

If the connect succeeds, the connection becomes the current one and ij displays a new prompt for the next command to be entered. If you have more than one open connection, the name of the connection appears in the prompt.

All further commands are processed against the new, current connection.

Example

```
ij> connect 'jdbc:derby:menuDB;create=true';
ij> -- we create a new table in menuDB:
CREATE TABLE menu(course CHAR(10), item CHAR(20), price INTEGER);
ij> protocol 'jdbc:derby:';
ij> connect 'sample' as sample1;
ij(SAMPLE1)> connect 'newDB;create=true' as newDB;
ij(NEWDB)> show connections;
CONNECTION0 - jdbc:derby:menuDB
NEWDB* - jdbc:derby:anotherDB
SAMPLE1 - jdbc:derby:newDB
ij>
ij> connect 'jdbc:derby:sample' user 'sa' password 'cloud3x9';
ij>
```

Describe

Syntax

```
DESCRIBE { table-Name | view-Name }
```

Description

Provides a description of the specified table or view. For a list of tables in the current schema, use the [Show Tables](#) command. For a list of views in the current schema, use the [Show Views](#) command. For a list of available schemas, use the [Show Schemas](#) command.

Example

```
ij> describe airlines;
```

COLUMN_NAME	TYPE_NAME	DEC	NUM	COLUM	COLUMN_DEF	CHAR_OCTE	IS_NULL
AIRLINE	CHAR			2		NULL	NO
AIRLINE_FULL	VARCHAR			24		NULL	YES
BASIC_RATE	DOUBLE		2	52		NULL	YES
DISTANCE_DISCOUNT	DOUBLE		2	52		NULL	YES
BUSINESS_LEVEL_FACT	DOUBLE		2	52		NULL	YES
FIRSTCLASS_LEVEL_FA	DOUBLE		2	52		NULL	YES
ECONOMY_SEATS	INTEGER	0	10	10		NULL	YES
BUSINESS_SEATS	INTEGER	0	10	10		NULL	YES
FIRSTCLASS_SEATS	INTEGER	0	10	10		NULL	YES

Disconnect

Syntax

```
DISCONNECT [ ALL | CURRENT | ConnectionIdentifier ]
```

Description

Disconnects from the database. Specifically issues a `java.sql.Connection.close` request against the connection indicated on the command line. There must be a current connection at the time the request is made.

If ALL is specified, all known connections are closed and there will be no current connection.

Disconnect CURRENT is the same as Disconnect without indicating a connection, the default connection is closed.

If a connection name is specified with an identifier, the command disconnects the named connection. The name must be the name of a connection in the current session provided with the `ij.connection.connectionName` property or with the `Connect` command.

If the `ij.database` property or the `Connect` command without the AS clause was used, you can supply the name the system generated for the connection. If the current connection is the named connection, when the command completes, there will be no current connection and you must issue a `Set Connection` or `Connect` command.

A Disconnect command issued against a Derby connection does not shut down the database or Derby (but the `Exit` command does).

Example

```
ij> connect 'jdbc:derby:menuDB;create=true';
ij> -- we create a new table in menuDB:
CREATE TABLE menu(course CHAR(10), ITEM char(20), PRICE integer);
0 rows inserted/updated/deleted
ij> disconnect;

ij> protocol 'jdbc:derby:';
ij> connect 'sample' as sample1;
ij> connect 'newDB;create=true' as newDB;
SAMPLE1 - jdbc:derby:sample
NEWDB* - jdbc:derby:newDB;create=true
* = current connection
ij(NEWDB)> set connection sample1;
ij> disconnect sample1;
ij> disconnect all;
ij>
```

Driver

Syntax

DRIVER *DriverNameString*

Description

Takes the value of the *DriverNameString* and issues a *Class.forName* request to load the named class. The class is expected to be a JDBC driver that registers itself with *java.sql.DriverManager*.

If the **Driver** command succeeds, a new *ij* prompt appears for the next command.

Example

```
ij> -- load the Derby driver so that a connection
-- can be made:
driver 'org.apache.derby.jdbc.EmbeddedDriver';
ij> connect 'jdbc:derby:menuDB;create=true';
ij>
```

Elapsedtime

Syntax

ELAPSED TIME { ON | OFF }

Description

When *elapsedtime* is turned on, *ij* displays the total time elapsed during statement execution. The default value is OFF.

Example

```
ij> elapsedtime on;
ij> VALUES current_date;
1
-----
1998-07-15
ELAPSED TIME = 2134 milliseconds
ij>
```

Execute

Syntax

EXECUTE { *SQLString* | *PreparedStatementIdentifier* }
[USING { *String* | *Identifier* }]

Description

Has several uses:

- To execute an SQL command entered as *SQLString*, using the **ExecuteString** style. The String is passed to the connection without further examination or processing by *ij*. *Normally, you execute SQL commands directly, not with the Execute command.*
- To execute a named command identified by *PreparedStatementIdentifier*. This must be previously prepared with the *ij* **Prepare** command.
- To execute either flavor of command when that command contains dynamic parameters, specify the values in the Using portion of the command. In this style, the *SQLString* or previously prepared *PreparedStatementIdentifier* is executed using the values supplied as *String* or *Identifier*. The *Identifier* in the USING clause must have a result set as its result. Each row of the result set is applied to the input parameters of the command to be executed, so the number of columns in the Using's result set must match the number of input parameters in the **Execute**'s statement. The results of each execution of the **Execute** statement are displayed as they are made. If the Using's result set contains no rows, the **Execute**'s statement is

not executed.

When auto-commit mode is on, the Using's result set is closed upon the first execution of the [Execute](#) statement. To ensure multiple-row execution of the [Execute](#) command, use the [Autocommit](#) command to turn auto-commit off.

Example

```
ij> autocommit off;
ij> prepare menuInsert as 'INSERT INTO menu VALUES (?, ?, ?)';
ij> execute menuInsert using 'VALUES
('entree', 'lamb chop', 14),
('dessert', 'creme brulee', 6)';
1 row inserted/updated/deleted
1 row inserted/updated/deleted
ij> commit;
```

Exit

Syntax

EXIT

Description

Causes the `ij` application to complete and processing to halt. Issuing this command from within a file started with the [Run](#) command or on the command line causes the outermost input loop to halt.

`ij` automatically shuts down a Derby database running in an embedded environment (issues a `Connect 'jdbc:derby::shutdown=true'` request) on exit.

`ij` exits when the [Exit](#) command is entered or if given a command file on the Java invocation line, when the end of the command file is reached.

Example

```
ij> disconnect;
ij> exit;
C:\>
```

First

Syntax

FIRST *Identifier*

Description

Moves the cursor to the first row in the *ResultSet*, then fetches the row. The cursor must have been created with the [Get Scroll Insensitive Cursor](#) command. It displays a banner and the values of the row.

Note: This command works only in a Java 2 Platform, Standard Edition, v 1.2 (J2SE) or higher environment.

Example

```
ij> get scroll insensitive cursor scrollCursor as
'SELECT * FROM menu FOR UPDATE OF price';
ij> first scrollcursor;
COURSE      | ITEM                | PRICE
-----
entree      | lamb chop           | 14
```

Get Cursor

Syntax

```
GET [WITH {HOLD|NOHOLD}] CURSOR Identifier AS String
```

WITH HOLD is the default attribute of the cursor. For a non-holdable cursor, use the WITH NOHOLD option.

Note: WITH NOHOLD is only available in Java 2 Platform, Standard Edition, v 1.4 (J2SE) or higher.

Description

Creates a cursor with the name of the *Identifier* by issuing a *java.sql.Statement.executeQuery* request on the value of the *String*.

If the *String* is a statement that does not generate a result set, the behavior of the underlying database determines whether an empty result set or an error is issued. If there is an error in executing the statement, no cursor is created.

ij sets the cursor name using a *java.sql.Statement.setCursorName* request. Behavior with respect to duplicate cursor names is controlled by the underlying database. Derby does not allow multiple open cursors with the same name.

Once a cursor has been created, the *ijNext* and *Close* commands can be used to step through its rows, and if the connection supports positioned update and delete commands, they can be issued to alter the rows.

Example

```
ij> -- autocommit needs to be off so that the positioned update
ij> -- can see the cursor it operates against.
ij> autocommit off;
ij> get cursor menuCursor as
'SELECT * FROM menu FOR UPDATE OF price';
ij> next menuCursor;
COURSE      |ITEM                                |PRICE
-----
entree      |lamb chop                          |14
ij> next menuCursor;
COURSE      |ITEM                                |PRICE
-----
dessert     |creme brulee                       |6
ij> UPDATE menu SET price=price+1 WHERE CURRENT OF menuCursor;
1 row inserted/updated/deleted
ij> next menuCursor;
COURSE      |ITEM                                |PRICE
-----
appetizer   |baby greens salad                  |7
ij> close menuCursor;
ij> commit;
ij>
```

Get Scroll Insensitive Cursor

Syntax

```
GET SCROLL INSENSITIVE [WITH {HOLD|NOHOLD}]
CURSOR Identifier AS
String
```

WITH HOLD is the default attribute of the cursor. For a non-holdable cursor, use the WITH NOHOLD option.

Note: WITH NOHOLD is only available in Java 2 Platform, Standard Edition, v 1.4 (J2SE) or higher.

Description

Creates a scrollable insensitive cursor with the name of the *Identifier*. (It does this by issuing a `createStatement(ResultSet.TYPE_SCROLL_INSENSITIVE, ResultSet.CONCUR_READ_ONLY)` call and then executing the statement with `java.sql.Statement.executeQuery` request on the value of the String.)

Note: This command only works in a Java 2 Platform, Standard Edition, v 1.2 (J2SE) or higher environment.

If the *String* is a statement that does not generate a result set, the behavior of the underlying database determines whether an empty result set or an error is issued. If there is an error in executing the statement, no cursor is created.

ij sets the cursor name using a `java.sql.Statement.setCursorName` request. Behavior with respect to duplicate cursor names is controlled by the underlying database. Derby does not allow multiple open cursors with the same name.

Once a scrollable cursor has been created, you can use the follow commands to work with the result set:

- [Absolute](#)
- [After Last](#)
- [Before First](#)
- [Close](#)
- [First](#)
- [Last](#)
- [Next](#)
- [Previous](#)
- [Relative](#)

Example

```
ij> autocommit off;
ij> get scroll insensitive cursor scrollCursor as
'SELECT * FROM menu';
ij> absolute 5 scrollCursor;
COURSE      | ITEM                      | PRICE
-----
entree      | lamb chop                 | 14
ij> after last scrollCursor;
No current row
ij> before first scrollCursor;
No current row
ij> first scrollCursor;
COURSE      | ITEM                      | PRICE
-----
entree      | lamb chop                 | 14
ij> last scrollCursor;
COURSE      | ITEM                      | PRICE
-----
dessert     | creme brulee              | 6
ij> previous scrollCursor;
COURSE      | ITEM                      | PRICE
-----
entree      | lamb chop                 | 14
ij> relative 1 scrollCursor;
COURSE      | ITEM                      | PRICE
-----
dessert     | creme brulee              | 6
ij>>previous scrollCursor;
COURSE      | ITEM                      | PRICE
-----
dessert     | creme brulee              | 6
ij> next scrollCursor;
COURSE      | ITEM                      | PRICE
-----
dessert     | creme brulee              | 6
```

Help

Syntax

HELP**Description**

Prints out a brief list of the `ij` commands.

Last**Syntax**

`LAST Identifier`

Description

Moves the cursor to the last row in the *ResultSet*, then fetches the row. The cursor must have been created with the [Get Scroll Insensitive Cursor](#) command. It displays a banner and the values of the row.

Note: This command works only in a Java 2 Platform, Standard Edition, v 1.2 (J2SE) or higher environment.

Example

```
ij> get scroll insensitive cursor scrollCursor as
'SELECT * FROM menu FOR UPDATE OF price';
ij> absolute 5 scrollCursor;
COURSE      | ITEM                | PRICE
-----|-----
entree      | lamb chop           | 14
ij> last scrollCursor;
COURSE      | ITEM                | PRICE
-----|-----
dessert     | creme brulee        | 6
```

LocalizedDisplay**Syntax**

`LOCALIZEDDISPLAY { on | off }`

Description

Specifies to display locale-sensitive data (such as dates) in the native format for the `ij` locale. The `ij` locale is the same as the Java system locale.

Note: NUMERIC and DECIMAL values are not localized when using the J2ME/CDC/Foundation Profile because of platform limitations.

Example

The following demonstrates *LocalizedDisplay* in an English locale:

```
ij> VALUES CURRENT_DATE;
1
-----
2000-05-01
1 row selected
ij> localizeddisplay on;
ij> VALUES CURRENT_DATE;
1
-----
May 1, 2000
1 row selected
```

MaximumDisplayWidth**Syntax**

MAXIMUMDISPLAYWIDTH *integer_value*

Description

Sets the largest display width for columns to the specified value. This is generally used to increase the default value in order to display large blocks of text.

Example

```
ij> maximumdisplaywidth 3;
ij> VALUES 'NOW IS THE TIME!';
1
---
NOW
ij> maximumdisplaywidth 30;
ij> VALUES 'NOW IS THE TIME!';
1
-----
NOW IS THE TIME!
```

Next

Syntax

NEXT *Identifier*

Description

Fetches the next row from the named cursor created with the [Get Cursor](#) command or [Get Scroll Insensitive Cursor](#). It displays a banner and the values of the row.

Example

```
ij> get cursor menuCursor as 'SELECT * FROM menu';
ij> next menuCursor;
COURSE      |ITEM                        |PRICE
-----
entree      |lamb chop                  |14
ij>
```

Prepare

Syntax

PREPARE *Identifier* **AS** *String*

Description

Creates a *java.sql.PreparedStatement* using the value of the String, accessible in *ij* by the *Identifier* given to it. If a prepared statement with that name already exists in *ij*, an error will be returned and the previous prepared statement will remain. Use the [Remove](#) command to remove the previous statement first. If there are any errors in preparing the statement, no prepared statement is created.

Any SQL statements allowed in the underlying connection's prepared statement can be prepared with this command.

Example

```
ij> prepare seeMenu as 'SELECT * FROM menu';
ij> execute seeMenu;
COURSE      |ITEM                        |PRICE
-----
entree      |lamb chop                  |14
dessert     |creme brulee               |6
```

```
2 rows selected
ij>
```

Previous

Syntax

```
PREVIOUS Identifier
```

Description

Moves the cursor to the row previous to the current one, then fetches the row. The cursor must have been created with the [Get Scroll Insensitive Cursor](#) command. It displays a banner and the values of the row.

Note: This command works only in a Java 2 Platform, Standard Edition, v 1.2 (J2SE) or higher environment.

Example

```
ij> get scroll insensitive cursor scrollCursor as
'SELECT * FROM menu FOR UPDATE OF price';
ij> last scrollcursor;
COURSE      | ITEM                | PRICE
-----
dessert     | creme brulee        | 6
ij> previous scrollcursor;
COURSE      | ITEM                | PRICE
-----
entree      | lamb chop           | 14
```

Protocol

Syntax

```
PROTOCOL String [ AS Identifier ]
```

Description

Specifies the protocol, as a String, for establishing connections and automatically loads the appropriate driver. *Protocol* is the part of the database connection URL syntax appropriate for your environment, including the JDBC protocol and the protocol specific to Derby. For further information about the Derby database connection URL, see the *Derby Developer's Guide*. Only Derby protocols are supported. Those protocols are listed in [ij.protocol](#).

Providing a protocol allows you to use a shortened database connection URL for connections. You can provide only the database name instead of the full protocol. In addition, you do not need to use the [Driver](#) command or specify a driver at start-up, since the driver is loaded automatically.

If you name the protocol, you can refer to the protocol name in the [Connect](#) command.

Example

```
ij> protocol 'jdbc:derby:';
ij> connect 'sample';
```

Readonly

Syntax

```
READONLY { ON | OFF }
```

Description

Sets the current connection to a "read-only" connection, as if the current user were defined as a *readOnlyAccess* user. (For more information about database authorization, see the *Derby Developer's Guide*.)

Example

```
ij> connect 'jdbc:derby:menuDB';
ij> readonly on;
ij> SELECT * FROM menu;
COURSE      | ITEM                      | PRICE
-----|-----|-----
entree      | lamb chop                 | 14
dessert     | creme brulee              | 6
appetizer   | baby greens               | 7
entree      | lamb chop                 | 14
entree      | lamb chop                 | 14
dessert     | creme brulee              | 6
6 rows selected
ij> UPDATE menu set price = 3;
ERROR 25502: An SQL data change is not permitted for a read-only
connection, user or database.
```

Relative

Syntax

```
RELATIVE int Identifier
```

Description

Moves the cursor to the row that is *int* number of rows relative to the current row, then fetches the row. The cursor must have been created with the [Get Scroll Insensitive Cursor](#) command. It displays a banner and the values of the row.

Note: This command works only in a Java 2 Platform, Standard Edition, v 1.2 (J2SE) or higher environment.

Example

```
ij> -- autocommit needs to be off so that the positioned update
ij> -- can see the cursor it operates against.
ij> autocommit off;
ij> get scroll insensitive cursor scrollCursor as
'SELECT * FROM menu FOR UPDATE OF price';
ij> last scrollcursor;
COURSE      | ITEM                      | PRICE
-----|-----|-----
dessert     | creme brulee              | 6
ij> previous scrollcursor;
COURSE      | ITEM                      | PRICE
-----|-----|-----
entree      | lamb chop                 | 14
ij> relative 1 scrollcursor;
COURSE      | ITEM                      | PRICE
-----|-----|-----
dessert     | creme brulee              | 6
```

Remove

Syntax

```
REMOVE Identifier
```

Description

Removes a previously prepared statement from ij. The identifier is the name by which the statement was prepared. The statement is closed to release its database resources.

Example

```
ij> prepare seeMenu as 'SELECT * FROM menu';
```

```

ij> execute seeMenu;
COURSE      | ITEM                      | PRICE
-----|-----|-----
entree      | lamb chop                 | 14
dessert     | creme brulee              | 6

2 rows selected
ij> remove seeMenu;
ij> execute seeMenu;
IJ ERROR: Unable to establish prepared statement SEEMENU
ij>

```

Rollback

Syntax

ROLLBACK

Description

Issues a *java.sql.Connection.rollback* request. Use only if auto-commit is off. A *java.sql.Connection.rollback* request undoes the currently active transaction and initiates a new transaction.

Example

```

ij> autocommit off;
ij> INSERT INTO menu VALUES ('dessert', 'rhubarb pie', 4);
1 row inserted/updated/deleted
ij> SELECT * from menu;
COURSE      | ITEM                      | PRICE
-----|-----|-----
entree      | lamb chop                 | 14
dessert     | creme brulee              | 7
appetizer   | baby greens               | 7
dessert     | rhubarb pie               | 4

4 rows selected
ij> rollback;
ij> SELECT * FROM menu;
COURSE      | ITEM                      | PRICE
-----|-----|-----
entree      | lamb chop                 | 14
dessert     | creme brulee              | 7
appetizer   | baby greens               | 7

3 rows selected
ij>

```

Run

Syntax

RUN *String*

Description

Assumes that the value of the string is a valid file name, and redirects *ij* processing to read from that file until it ends or an [Exit](#) command is executed. If the end of the file is reached without *ij* exiting, reading will continue from the previous input source once the end of the file is reached. Files can contain Run commands.

ij prints out the statements in the file as it executes them.

Any changes made to the *ij* environment by the file are visible in the environment when processing resumes.

Example

```

ij> run 'setupMenuConn.ij';

```

```

ij> -- this is setupMenuConn.ij
-- ij displays its contents as it processes file
ij> connect 'jdbc:derby:menuDB';
ij> autocommit off;
ij> -- this is the end of setupMenuConn.ij
-- there is now a connection to menuDB and no autocommit.
-- input will now resume from the previous source.
;
ij>

```

Set Connection

Syntax

```
SET CONNECTION Identifier
```

Description

Allows you to specify which connection to make current when you have more than one connection open. Use the [Show Connections](#) command to display open connections.

If there is no such connection, an error results and the current connection is unchanged.

Example

```

ij> protocol 'jdbc:derby:';
ij> connect 'sample' as sample1;
ij> connect 'newDB;create=true' as newDB;
ij (NEWDB)> show connections;
SAMPLE1 -      jdbc:derby:sample
NEWDB* -      jdbc:derby:newDB;create=true
* = current connection
ij (NEWDB)> set connection sample1;
ij (SAMPLE1)> disconnect all;
ij>

```

Show

Syntax

```

SHOW
{
  CONNECTIONS |
  INDEXES [ IN schemaName | FROM table-Name ] |
  PROCEDURES [ IN schemaName ] |
  SCHEMAS |
  SYNONYMS [ IN schemaName ] |
  TABLES [ IN schemaName ] |
  VIEWS [ IN schemaName ] |
}

```

Description

The SHOW command can be used to display information about active connections and database objects.

SHOW CONNECTIONS

If there are no connections, the SHOW CONNECTIONS command returns "No connections available".

Otherwise, the command displays a list of connection names and the URLs used to connect to them. The currently active connection, if there is one, is marked with an * after its name.

Example

```

ij> connect 'sample' as sample1;
ij> connect 'newDB;create=true' as newDB;
ij(NEWDB)> show connections;
SAMPLE1 -          jdbc:derby:sample
NEWDB* -          jdbc:derby:newDB;create=true
* = current connection
ij(NEWDB)>

```

SHOW INDEXES

SHOW INDEXES will display all the indexes in the database.

If **IN *schemaName*** is specified, then only the indexes in the specified schema are displayed.

If **FROM *table-Name*** is specified, then only the indexes on the specified table are displayed.

Example

```

ij> show indexes in app;

```

TABLE_NAME	COLUMN_NAME	NON_U&	TYPE	ASC&	CARDINA&	PAGES
AIRLINES	AIRLINE	false	3	A	NULL	NULL
COUNTRIES	COUNTRY_ISO_CODE	false	3	A	NULL	NULL
COUNTRIES	COUNTRY	false	3	A	NULL	NULL
CITIES	CITY_ID	false	3	A	NULL	NULL
FLIGHTS	FLIGHT_ID	false	3	A	NULL	NULL
FLIGHTS	SEGMENT_NUMBER	false	3	A	NULL	NULL
FLIGHTAVAILABILITY	FLIGHT_ID	false	3	A	NULL	NULL
FLIGHTAVAILABILITY	SEGMENT_NUMBER	false	3	A	NULL	NULL
FLIGHTAVAILABILITY	FLIGHT_DATE	false	3	A	NULL	NULL
MAPS	MAP_ID	false	3	A	NULL	NULL
MAPS	MAP_NAME	false	3	A	NULL	NULL
FLIGHTS	DEST_AIRPORT	true	3	A	NULL	NULL
FLIGHTS	ORIG_AIRPORT	true	3	A	NULL	NULL
CITIES	COUNTRY_ISO_CODE	true	3	A	NULL	NULL
FLIGHTAVAILABILITY	FLIGHT_ID	true	3	A	NULL	NULL
FLIGHTAVAILABILITY	SEGMENT_NUMBER	true	3	A	NULL	NULL

16 rows selected

```

ij> show indexes from flights;

```

TABLE_NAME	COLUMN_NAME	NON_U&	TYPE	ASC&	CARDINA&	PAGES
FLIGHTS	FLIGHT_ID	false	3	A	NULL	NULL
FLIGHTS	SEGMENT_NUMBER	false	3	A	NULL	NULL
FLIGHTS	DEST_AIRPORT	true	3	A	NULL	NULL
FLIGHTS	ORIG_AIRPORT	true	3	A	NULL	NULL

4 rows selected

SHOW PROCEDURES

SHOW PROCEDURES will display all the procedures in the database that have been created with the CREATE PROCEDURE statement, as well as system procedures.

If **IN *schemaName*** is specified, only procedures in the specified schema are displayed.

Example

```

ij> show procedures in syscs_util;

```

PROCEDURE_SCHEM	PROCEDURE_NAME	REMARKS
SYSCS_UTIL	SYSCS_BACKUP_DATABASE	org.apache.derby.ca&
SYSCS_UTIL	SYSCS_BACKUP_DATABASE_AND_ENA&	org.apache.derby.ca&
SYSCS_UTIL	SYSCS_BACKUP_DATABASE_AND_ENA&	org.apache.derby.ca&
SYSCS_UTIL	SYSCS_BACKUP_DATABASE_NOWAIT	org.apache.derby.ca&
SYSCS_UTIL	SYSCS_BULK_INSERT	org.apache.derby.ca&
SYSCS_UTIL	SYSCS_CHECKPOINT_DATABASE	org.apache.derby.ca&
SYSCS_UTIL	SYSCS_COMPRESS_TABLE	org.apache.derby.ca&
SYSCS_UTIL	SYSCS_DISABLE_LOG_ARCHIVE_MODE	org.apache.derby.ca&
SYSCS_UTIL	SYSCS_EXPORT_QUERY	org.apache.derby.ca&

SYSCS_UTIL	SYSCS_EXPORT_TABLE	org.apache.derby.ca&
SYSCS_UTIL	SYSCS_FREEZE_DATABASE	org.apache.derby.ca&
SYSCS_UTIL	SYSCS_IMPORT_DATA	org.apache.derby.ca&
SYSCS_UTIL	SYSCS_IMPORT_TABLE	org.apache.derby.ca&
SYSCS_UTIL	SYSCS_INPLACE_COMPRESS_TABLE	org.apache.derby.ca&
SYSCS_UTIL	SYSCS_SET_DATABASE_PROPERTY	org.apache.derby.ca&
SYSCS_UTIL	SYSCS_SET_RUNTIMESTATISTICS	org.apache.derby.ca&
SYSCS_UTIL	SYSCS_SET_STATISTICS_TIMING	org.apache.derby.ca&
SYSCS_UTIL	SYSCS_UNFREEZE_DATABASE	org.apache.derby.ca&

18 rows selected

SHOW SCHEMAS

SHOW SCHEMAS will display all of the schemas in the current connection.

Example

```
ij> show schemas;
TABLE_SCHEM
```

```
-----
APP
NULLID
SQLJ
SYS
SYSCAT
SYSCS_DIAG
SYSCS_UTIL
SYSFUN
SYSIBM
SYSPROC
SYSSTAT
```

11 rows selected

SHOW SYNONYMS

SHOW SYNONYMS will display all the synonyms in the database that have been created with the CREATE SYNONYMS statement.

If IN *schemaName* is specified, only synonyms in the specified schema are displayed.

Example

```
ij> show synonyms;
```

TABLE_SCHEM	TABLE_NAME	REMARKS
APP	MYAIRLINES	

SHOW TABLES

SHOW TABLES will display all of the tables in the current schema.

If IN *schemaName* is specified, the tables in the given schema are displayed.

Example

```
ij> show tables;
```

TABLE_SCHEM	TABLE_NAME	REMARKS
APP	AIRLINES	
APP	CITIES	
APP	COUNTRIES	
APP	FLIGHTAVAILABILITY	
APP	FLIGHTS	
APP	FLIGHTS_HISTORY	
APP	MAPS	

7 rows selected

SHOW VIEWS

SHOW VIEWS will display all of the views in the current schema.

If **IN *schemaName*** is specified, the tables in the given schema are displayed.

Example

```
ij> show views;
TABLE_SCHEM | TABLE_NAME | REMARKS
-----|-----|-----
APP | TOTALSEATS |
1 row selected
```

Wait For

Syntax

```
WAIT FOR Identifier
```

Description

Displays the results of a previously started asynchronous command.

The identifier for the asynchronous command must have been used in a previous [Async](#) command on this connection. The [Wait For](#) command waits for the SQL statement to complete execution, if it has not already, and then displays the results. If the statement returns a result set, the [Wait For](#) command steps through the rows, not the [Async](#) command. This might result in further execution time passing during the result display.

Example

See [Async](#).

Comment

Syntax

```
-- Text
```

Description

You can use a double dash to create a comment anywhere within an `ij` command line and as permitted by the underlying connection within SQL commands. The comment is ended at the first new line encountered in the text.

Comments are ignored on input and have no effect on the output displayed.

Example

```
ij> -- this is a comment;
-- the semicolons in the comment are not taken as the end
-- of the command; for that, we put it outside the --:
;
ij>
```

Identifier

Syntax

Identifier**Description**

Some `ij` commands require identifiers. These `ij` identifiers are case-insensitive. They must begin with a letter in the range A-Z, and can consist of any number of letters in the range A-Z, digits in the range 0-9, and underscore (`_`) characters.

These identifiers exist within the scope of `ij` only and are distinct from any identifiers used in SQL commands, except in the case of the [Get Cursor](#) command. The [Get Cursor](#) command specifies a cursor name to use in creating a result set.

`ij` does not recognize or permit delimited identifiers in `ij` commands. They can be used in SQL commands.

Example

```
These are valid ij identifiers:
fool
exampleIdentifier12345
another_one
```

String**Syntax**

```
'Text'
```

Description

Some `ij` commands require strings. `ij` strings are represented by the same literal format as SQL strings and are delimited by single quotation marks. To include a single quotation mark in a string, you must use two single quotation marks, as shown in the examples below. `ij` places no limitation on the lengths of strings, and will treat embedded new lines in the string as characters in the string.

Some `ij` commands execute SQL commands specified as strings. Therefore, you must double any single quotation marks within such strings, as shown in the second example below.

The cases of letters within a string are preserved.

Example

```
This is a string in ij      And this is its value
'Mary's umbrella'         Mary's umbrella
'hello world'              hello world

--returns Joe's
execute 'VALUES ''Joe''''s''';
```

ij errors

`ij` might issue messages to inform the user of errors during processing of statements.

ERROR SQLState

When the underlying JDBC driver returns an *SQLException*, `ij` displays the *SQLException* message with the prefix "ERROR SQLState". If the *SQLException* has no SQLState associated with it, the prefix "ERROR (no SQLState)" is used.

WARNING SQLState

Upon completion of execution of any JDBC request, *ij* will issue a *getWarnings* request and display the SQLWarnings that are returned. Each *SQLWarning* message is displayed with the prefix "WARNING SQLState". If an *SQLWarning* has no SQLState associated with it, the prefix "WARNING (no SQLState)" is used.

IJ ERROR

When *ij* runs into errors processing user commands, such as being unable to open the file named in a [Run](#) command or not having a connection to disconnect from, it prints out a message with the prefix "IJ ERROR".

IJ WARNING

ij displays warning messages to let the user know if behavior might be unexpected. *ij* warnings are prefixed with "IJ WARNING".

JAVA ERROR

When an unexpected Java exception occurs, *ij* prints a message with the prefix "JAVA ERROR".

Using the bulk import and export procedures

You might want to import or export a large amount of data between files and the database. Instead of having to use INSERT and SELECT statements, you can use Derby procedures to import data directly from files into tables and to export data from tables into files.

Bulk Import/Export overview

Derby provides import and export system procedures that you can use to import and export data in delimited data file format.

- Use export procedures to write data from a database to one or more files that are stored outside of the database. You can use a procedure to export data from a table into a file or export data from a SELECT statement result into a file.
- Use import procedures to import data from a file into a table. If the target table already contains data, you can replace or append to the existing data.

You can perform an Import or Export operation from `ij` or from within an SQL statement.

Options for running the import and export procedures

You can run the import/export procedures from within an SQL statement using `ij` or any Java application.

Import/Export reads and writes only text files. Import does not support read-once streams (live data feeds), because it reads the first line of the file to determine the number of columns, then reads it again to import the data.

Note: These server-side utilities exhibit different behavior in client/server mode. Typically, you use them to import data into and export data from a locally running Derby. However, you can use the import/export procedures when Derby is running in a server framework if you specify import and export files that are accessible to the server.

Bulk import/export requirements and considerations

The table must exist

For you to import data into a table, the table must already exist in Derby. The table does not have to be empty. If the table is not empty, bulk import performs single inserts which results in slower performance.

Create indexes and primary key, foreign key, and unique constraints first

To avoid a separate create index step, create indexes and primary keys on tables before you import data. However, if your memory and disk spaces resources are limited, you can build the indexes and primary keys after importing data.

Data types

You can import and export only data of the non-binary, *built-in* data types. Derby implicitly converts the strings to the data type of the receiving column. If any of the implicit conversions fail, the whole import is aborted. For example, "3+7" cannot be converted into an integer. An export that encounters a runtime error stops.

Note: You cannot import or export the binary data types: BLOB, CLOB, CHAR FOR BIT DATA, VARCHAR FOR BIT DATA, or LONG VARCHAR FOR BIT DATA.

Additionally, you cannot import or export the XML data type.

Locking during import

Import procedures use the same isolation level as the connection in which they are executed to insert data into tables. During import, the entire table is exclusively locked irrespective of the isolation level.

Locking during export

Export procedures use the same isolation level as the connection in which they are executed to fetch data from tables.

Import behavior on tables with triggers

The import procedure fires INSERT triggers when data is appended to the table. The

REPLACE option is not allowed when triggers are enabled on the table.

Restrictions on the REPLACE option

If a table that receives imported data already contains data, you can either replace or append to the existing data. You can use the REPLACE option on tables that have dependent tables, but the replaced data must maintain referential integrity, otherwise the import operation will be rolled back. You cannot use the REPLACE option if the table has enabled triggers.

Restrictions on tables

You cannot use import procedures to import data into a system table or a declared temporary table.

Bulk import and export

You can use import and export procedures which are executable from ij or any Java program. You must have derbytools.jar in your classpath to use the import or export procedures from ij.

To invoke an import or export procedure, you must be connected to the database into which data is imported or from which data is exported. Because the procedures will issue a COMMIT or a ROLLBACK statement, you should perform either a COMMIT or ROLLBACK to complete all transactions and release all table-level locks prior to invoking the import or export procedure. Other user applications that access the table with a separate connection do not need to disconnect.

Note: Imports are transactional. If an error occurs during bulk import, all changes are rolled back.

Bulk-Import

Derby provides two import procedures you can use to perform bulk-import operations:

1. To import data from a file to a table, use the SYSCS_UTIL.SYSCS_IMPORT_TABLE procedure. The procedure definition is:

```
SYSCS_UTIL.SYSCS_IMPORT_TABLE (IN schemaName VARCHAR(128),
IN tableName VARCHAR(128), IN fileName VARCHAR(32672),
IN columnDelimiter CHAR(1), IN characterDelimiter CHAR(1),
IN codeset VARCHAR(128), IN replace SMALLINT)
```

No Result is returned from the procedure.

2. To import data from a file to a subset of columns in a table, use the SYSCS_UTIL.SYSCS_IMPORT_DATA procedure. To import data to subset of columns in a table, you specify *insertColumns* on the table into which data will be imported and/or specify *columnIndexes* to import data fields from a file to columns in a table. The procedure definition is:

```
SYSCS_UTIL.SYSCS_IMPORT_DATA (IN schemaName VARCHAR(128),
IN tableName VARCHAR(128), IN insertColumns VARCHAR(32672),
IN columnIndexes VARCHAR(32672), IN fileName VARCHAR(32672),
IN columnDelimiter CHAR(1), IN characterDelimiter CHAR(1),
IN codeset VARCHAR(128), IN replace SMALLINT)
```

No result is returned from the procedure.

Arguments to the import procedure

- *schemaName*

Specifies the schema of the table. You can pass a NULL value to use the default schema name.

- *tableName*

Specifies the table name of the table into which the data is to be imported. This table cannot be a system table or a declared temporary table. The string must exactly match case of the table name. Passing a null will result in an error.

- *insertColumns*

Specifies the comma separated column names of the table into which the data will be imported. You can pass a NULL value to import into all columns of the table.

- *columnIndexes*

Specifies the comma separated column indexes (numbered from one) of the input data fields that will be imported. You can pass a NULL value to use all input data fields in the file.

- *fileName*

Specifies the file that contains the data to be imported. If the path is omitted, the current working directory is used. The specified location of the file should refer to the server side location if using the Network Server. Passing a null will result in an error.

- *columnDelimiter*

Specifies a column delimiter. The specified character is used in place of a comma to signify the end of a column. You can pass a NULL value to use the default value of a comma.

- *characterDelimiter*

Specifies a character delimiter. The specified character is used in place of double quotation marks to enclose a character string. You can pass a NULL value to use the default value of a double quotation mark.

- *codeset*

Specifies the code set of the data in the input file. The code set name should be one of the Java-supported character encoding sets. Data is converted from the specified code set to the database code set (UTF-8). You can pass a NULL value to interpret the data file in the same code set as the JVM in which it is being executed.

- *replace*

A non-zero value for the `replace` parameter will import in REPLACE mode, while a zero value will import in INSERT mode. REPLACE mode deletes all existing data from the table by truncating the table and inserts the imported data. The table definition and the index definitions are not changed. You can only import with REPLACE mode if the table already exists. INSERT mode adds the imported data to the table without changing the existing table data. Passing a null value will result in an error.

If you create a schema, table, or column name as a non-delimited identifier, you must pass the name to the import procedure using all upper-case characters. If you created a schema, table, or column name as a delimited identifier, you must pass the name to the import procedure using the same case that was used when it was created.

Bulk-Export

Derby provides two export procedures you can use to perform bulk-export operations:

1. To export all the data from a table to a file, use the SYSCS_UTIL.SYSCS_EXPORT_TABLE procedure. The procedure definition is:

```
SYSCS_UTIL.SYSCS_EXPORT_TABLE (IN schemaName VARCHAR(128),
IN tableName VARCHAR(128), IN fileName VARCHAR(32672),
IN columnDelimiter CHAR(1), IN characterDelimiter CHAR(1),
IN codeset VARCHAR(128))
```

No Result is returned from the procedure.

2. To export the result of a SELECT statement to a file, use the SYSCS_UTIL.SYSCS_EXPORT_QUERY procedure. The procedure definition is:

```
SYSCS_UTIL.SYSCS_EXPORT_QUERY (IN selectStatement VARCHAR(32672),
```

```
IN tableName VARCHAR(128), IN fileName VARCHAR(32672),
IN columnDelimiter CHAR(1), IN characterDelimiter CHAR(1),
IN codeset VARCHAR(128))
```

No result is returned from the procedure.

Arguments to the export procedure

- *schemaName*
Specifies the schema of the table. You can pass a NULL value to use the default schema name.
- *selectStatement*
Specifies the SELECT statement query that returns the data to be exported. Passing a NULL value will result in an error.
- *tableName*
Specifies the table name of the table or view from which the data is to be exported. This table cannot be a system table or a declared temporary table. The string must exactly match the case of the table name. Passing a null will result in an error.
- *fileName*
Specifies the file to which the data is to be exported. If the path is omitted, the current working directory is used. If the name of a file that already exists is specified, the export utility overwrites the contents of the file; it does not append the information. The specified location of the file should refer to the server-side location if using the Network Server. Passing a null will result in an error.
- *columnDelimiter*
Specifies a column delimiter. The specified character is used in place of a comma to signify the end of a column. You can pass a NULL value to use the default value of a comma.
- *characterDelimiter*
Specifies a character delimiter. The specified character is used in place of double quotation marks to enclose a character string. You can pass a NULL value to use the default value of a double quotation mark.
- *codeset*
Specifies the code set of the data in the export file. The code set name should be one of the Java-supported character encoding sets. Data is converted from the database code page to the specified code page before writing to the file. You can pass a NULL value to write the data in the same code page as the JVM in which it is being executed.

If you create a schema, table, or column name as a non-delimited identifier, you must pass the name to the export procedure using all upper-case characters. If you created a schema or table name as a delimited identifier, you must pass the name to the export procedure using the same case that was used when it was created.

Examples of bulk import and export

All examples in this section are run using the `ij` utility.

The following example shows how to import data into the *staff* table in a sample database from the `myfile.del` file.

```
CALL SYSCS_UTIL.SYSCS_IMPORT_TABLE (null,'staff','myfile.del',null,null,
null,0);
```


The following example shows how to import data into the *staff* table in a sample database from a delimited data file `myfile.del`. This example defines the percentage character (%) as the string delimiter, and a semicolon as the column delimiter.

```
CALL SYSCS_UTIL.SYSCS_IMPORT_TABLE
(null,'staff','c:\output\myfile.del',';', '%',
null,0);
```

The following example shows how to export data from the *staff* table in a sample database to the file `myfile.del`.

```
CALL SYSCS_UTIL.SYSCS_EXPORT_TABLE
(null,'staff','myfile.del',null,null,null);
```

The following example shows how to export employee data in department 20 from the *staff* table in a sample database to the file `awards.del`.

```
CALL SYSCS_UTIL.SYSCS_EXPORT_QUERY ('select * from staff where dept=20',
'c:\output\awards.del',null,null,null);
```

The following example shows how to export data from the *staff* table to a delimited data file `myfile.del` with the percentage character (%) as the character delimiter, and a semicolon as the column delimiter from the *staff* table.

```
CALL SYSCS_UTIL.SYSCS_EXPORT_TABLE
(null,'staff','c:\output\myfile.del',';', '%',
null);
```

Importing into tables with identity columns

You can use the SYSCS_UTIL.SYSCS_IMPORT_DATA procedure to import data into a table that contains an identity column. If the identity column is defined as GENERATED ALWAYS, an identity value is generated for a table row whenever the corresponding row field in the input file does not contain a value for the identity column. When a corresponding row field in the input file already contains a value for the identity column, the row cannot be inserted into the table and the import operation will fail. To prevent such scenarios, the following examples show how to specify arguments in the SYSCS_UTIL.SYSCS_IMPORT_DATA procedure to ignore data for the identity column from the file, and/or omit the column name from the insert column list.

If the REPLACE option is used during import, Derby resets its internal counter of the last identity value for a column to the initial value defined for the identity column.

Consider the following table that contains an identity column, *c2*:

```
CREATE TABLE tab1 (c1 CHAR(30), c2 INT GENERATED ALWAYS AS IDENTITY, c3
REAL,
c4 CHAR(1))
```

- Suppose you want to import data into *tab1* from a file `myfile.del` that does not have identity column information and `myfile.del` contains three fields with the following data:

```
Robert,45.2,J
Mike,76.9,K
Leo,23.4,I
```

To import data from `myfile.del` into the `tab1` table, explicitly list the column names for `tab1` without the identity column `c2` and execute the `SYSCS_UTIL.SYSCS_IMPORT_DATA` procedure as follows:

```
CALL SYSCS_UTIL.SYSCS_IMPORT_DATA (NULL, 'TAB1', 'C1,C3,C4' , null,
'myfile.del',null, null,null,0)
```

- Suppose you want import data into `tab1` from a file `empfile.del` that also has identity column information and the file contains three fields with the following data:

```
Robert,1,45.2,J
Mike,2,23.4,I
Leo,3,23.4,I
```

To import data from `empfile.del` into the `tab1` table, explicitly specify an insert column list without the identity column `c2` and specify the column indexes without identity column data and execute the `SYSCS_UTIL.SYSCS_IMPORT_DATA` procedure as follows:

```
CALL SYSCS_UTIL.SYSCS_IMPORT_DATA (NULL, 'TAB1', 'C1,C3,C4' ,
'1,3,4',
'empfile.del',null, null,null,0)
```

Executing import/export procedures from JDBC

You can execute import and export procedures from a JDBC program. The following code fragment shows how you might call the `SYSCS_UTIL.SYSCS_EXPORT_TABLE` procedure from Java. In this example, the procedure exports the `staff` table data in the default schema to the `staff.dat` file, using a percentage (%) character to specify a column delimiter.

```
PreparedStatement ps=conn.prepareStatement("CALL
SYSCS_UTIL.SYSCS_EXPORT_TABLE
(?,?,?,?,,?)");
ps.setString(1,null);
ps.setString(2,"STAFF");
ps.setString(3,"staff.dat");
ps.setString(4,"%");
ps.setString(5,null);
ps.setString(6,null);
ps.execute();
```

File format for input and output

The default file format is a delimited text file with the following characteristics:

- Rows are separated by a new line.
- Fields are separated by a comma (,)
- Character-based fields are delimited with double quotes (")

Before performing import or export operations, you must ensure that the chosen delimiter character is not contained in the data to be imported or exported. If you chose a delimiter character that is part of the data to be imported/exported unexpected errors might occur.

The following restrictions apply to column and character delimiters:

- Delimiters are mutually exclusive
- A delimiter cannot be a line-feed character, a carriage return, or a blank space.
- The default decimal point (.) cannot be a character delimiter.

The record delimiter is assumed to be a new-line character. The record delimiter should not be used as any other delimiter.

Character delimiters are permitted with the character-based fields (CHAR, VARCHAR, and LONG VARCHAR) of a file during import, any pair of character delimiters found between the enclosing character delimiters is imported into the database. For example, suppose you have the following character string:

```
"what a "great" day!"
```

The preceding character string gets imported into the database as:

```
What a "great" day!
```

During export, the rule applies in reverse. For example, suppose you have the following character string:

```
"The boot has a 3" heel."
```

The preceding character string gets exported to a file as:

```
"The boot has a 3" "heel."
```

The following example file shows four rows and four columns in default file format:

```
1,abc,22,def
22,,,"a is a zero-length string, b is null"
13,"hello",454,"world"
4,b and c are both null,,
```

The export procedure outputs the following values:

```
1,"abc",22,"def"
22,,,"a is a zero-length string, b is null"
13,"hello",454,"world"
4,"b and c are both null",,,
```

Treatment of NULLS

In a delimited file, a NULL value is exported as an empty field. The following example shows the export of a four-column row in which the third column is NULL:

```
7,95,,Happy Birthday
```

Import works the same way; an empty field is imported as a NULL value.

CODESET values for import/export

Import and export procedures accept arguments to specify codeset values. You can specify the codeset (character encoding) for import and export procedures to override the system default. The following table contains a sample of character encoding supported by JDK 1.x. To review the complete list of character encodings, refer to your Java documentation.

Table 3. Sample character encodings

This table contains sample character encodings supported by JDK1.x.

Character Encoding	Explanation
8859_1	ISO Latin-1
8859_2	ISO Latin-2
8859_7	ISO Latin/Greek
Cp1257	Windows Baltic
Cp1258	Windows Vietnamese
Cp437	PC Original
EUCJIS	Japanese EUC
GB2312	GB2312-80 Simplified Chinese
JIS	JIS
KSC5601	KSC5601 Korean

Character Encoding	Explanation
MacCroatian	Macintosh Croatian
MacCyrillic	Macintosh Cyrillic
SJIS	PC and Windows Japanese
UTF-8	Standard UTF-8

The following example shows how to specify UTF-8 encoding to export from the *staff* table:

```
CALL SYSCS_UTIL.SYSCS_EXPORT_TABLE  
(NULL, 'STAFF', 'staff.dat', NULL, NULL, 'UTF-8')
```

The following example shows how to specify UTF-8 encoding to import from the *staff* table:

```
CALL SYSCS_UTIL.SYSCS_IMPORT_TABLE  
(NULL, 'STAFF', 'staff.dat', NULL, NULL, 'UTF-8', 0)
```

Storing jar files in a database

`SQLJ.install_jar`, `SQLJ.remove_jar`, and `SQLJ.replace_jar`, are a set of procedures in the SQL schema that allow you to store jar files in the database.

Your jar file has a *physical name* (the name you gave it when you created it) and a *Derby name* (the Derby identifier you give it when you load it into a particular schema). Its Derby name is an *SQL92Identifier*; it can be delimited and must be unique within a schema. A single schema can store more than one jar file.

Adding a Jar File

To add a jar file using SQL syntax:

```
CALL SQLJ.install_jar('jarFilePath', qualifiedJarName, 0)
```

- `jarFilePath`

The path and physical name of the jar file to add or use as a replacement. For example:

```
d:/todays_build/tours.jar
```

- `qualifiedJarName`

The Derby name of the jar file, qualified by the schema name. Two examples:

```
MYSHEMA.Sample1
```

```
-- a delimited identifier.  
MYSHEMA."Sample2"
```

Removing a jar file

To remove a jar file using SQL syntax:

```
CALL SQLJ.remove_jar (qualifiedJarName, 0)
```

Replacing a jar file

To replace a jar file using SQL syntax:

```
CALL SQLJ.replace_jar('jarFilePath', qualifiedJarName)
```

- `jarFilePath`

The path and physical name of the jar file to add or use as a replacement. For example:

```
d:/todays_build/tours.jar
```

- `qualifiedJarName`

The Derby name of the jar file, qualified by the schema name. Two examples:

```
MYSCHEMA.Sample1 -- a delimited identifier.
```

```
MYSCHEMA."Sample2"
```

Installing a jar example

- Complete SQL example for installing a jar:

```
CALL SQLJ.install_jar('d:\todays_build\tours.jar',  
'APP."ToursLogic!"', 0);
```

For more information about storing classes in a database, see the *Derby Developer's Guide*.

sysinfo

Use the sysinfo utility to display information about your Java environment and Derby (including version information). To use sysinfo, either *derby.jar* or *derbytools.jar* must be in your classpath.

```
java org.apache.derby.tools.sysinfo
```

sysinfo example

```
$ java org.apache.derby.tools.sysinfo
----- Java Information -----
Java Version: 1.4.2_07
Java Vendor: Sun Microsystems Inc.
Java home: c:\p4main\jdk142\jre
Java classpath: c:\Derby_10\lib\derby.jar;c:\Derby_10\lib\derbytools.jar;
c:\Derby_10\lib\derbyLocale_de_DE.jar;e:\Derby_10\lib\derbyLocale_es.jar;
e:\Derby_10\lib\derbyLocale_fr.jar;c:\Derby_10\lib\derbyLocale_it.jar;
e:\Derby_10\lib\derbyLocale_ja_JP.jar;e:\Derby_10\lib\derbyLocale_ko_KR.jar;
e:\Derby_10\lib\derbyLocale_pt_BR.jar;e:\Derby_10\lib\derbyLocale_zh_CN.jar;
e:\Derby_10\lib\derbyLocale_zh_TW.jar
OS name: Windows 2000
OS architecture: x86
OS version: 5.0
Java user name: user1
Java user home: C:\Documents and Settings\myhome
Java user dir: E:\p4main\sysrest\myrtst7
java.specification.name: Java Platform API Specification
java.specification.version: 1.4
----- Derby Information -----
JRE - JDBC: J2SE 1.4.2 - JDBC 3.0
[c:\Derby_10\lib\derby.jar] 10.2.1.0 - (190628)
[c:\Derby_10\lib\derbytools.jar] 10.2.1.0 - (190628)
----- Locale Information -----
Current Locale : [English/United States [en_US]]
Found support for locale: [de_DE]
version: 10.2.1.0 - (190628)
Found support for locale: [es]
version: 10.2.1.0 - (190628)
Found support for locale: [fr]
version: 10.2.1.0 - (190628)
Found support for locale: [it]
version: 10.2.1.0 - (190628)
Found support for locale: [ja_JP]
version: 10.2.1.0 - (190628)
Found support for locale: [ko_KR]
version: 10.2.1.0 - (190628)
Found support for locale: [pt_BR]
version: 10.2.1.0 - (190628)
Found support for locale: [zh_CN]
version: 10.2.1.0 - (190628)
Found support for locale: [zh_TW]
```

When requesting help from Derby technical support or posting on the forum, include a copy of the information provided by sysinfo.

Using sysinfo to check the classpath

sysinfo provides an argument (*-cp*) which can be used to test the classpath.

```
java org.apache.derby.tools.sysinfo -cp
[ [ embedded ] [ server ] [ client ] [ tools ] [ anyClass.class ] ]
```

If your environment is set up correctly, the utility shows output indicating success.

You can provide optional arguments with *-cp* to test different environments. Optional arguments to *-cp* are:

- embedded
- server

- client
- tools
- *classname.class*

If something is missing from your classpath, the utility indicates what is missing. For example, if you neglected to include the directory containing the class named *SimpleApp* to your classpath, the utility would indicate this when the following command line was issued (type all on one line):

```
$ java org.apache.derby.tools.sysinfo -cp embedded SimpleApp.class
FOUND IN CLASS PATH:

Derby embedded engine library (derby.jar)

NOT FOUND IN CLASS PATH:

user-specified class (SimpleApp)
(SimpleApp not found.)
```


dblook

Use the `dblook` utility to view all or parts of the Data Definition Language (DDL) for a given database. You must place the `Derbyderbytools.jar` file in the classpath directory to use the `dblook` utility.

Using dblook

The syntax for the command to launch the `dblook` utility is:

```
java org.apache.derby.tools.dblook -d <databaseURL> [OPTIONS]
```

The value for *databaseUrl* is the complete URL for the database. Where appropriate, the URL includes any connection attributes or properties that might be required to access the database.

For example, to connect to the database 'myDB', the URL would simply be `'jdbc:derby:myDB'`; to connect using the Network Server to a database `'C:\private\tmp\myDB'` on a remote server (port 1527), the URL would be:

```
'jdbc:derby://localhost:1527/"C:\private\tmp\myDB"
;user=someusr;password=somepwd'
```

As with other Derby utilities, you must ensure that no other JVMs are started against the database when you call the `dblook` utility, or an exception will occur and will print to the `dblook.log`. If this exception is thrown, the `dblook` utility will quit. To recover, you must ensure that no other Derby applications running in a separate JVM are connected to the source database. These connections need to be shutdown. Once all existing JVMs running against the database have been shutdown, the `dblook` utility will execute successfully. You can also start the Derby Network server and run the `dblook` utility as a client application while other clients are connected to the server.

dblook options

The `dblook` utility options include:

-z <schemaName>

specifies the schema to which the DDL should be restricted. Only objects with the specified schema are included in the DDL file.

-t <tableOne> <tableTwo> ...

specifies the tables to which the DDL should be restricted. All tables with a name from this list will be included in the DDL file subject to `-z` limitations, as will the DDL for any keys, checks, or indexes on which the table definitions depend.

Additionally, if the statement text of any triggers or views includes a reference to any of the listed table names, the DDL for that trigger/view will also be generated, subject to `-z` limitations. If a table is not included in this list, then neither the table nor any of its keys, checks, or indexes will be included in the final DDL. If this parameter is not provided, all database objects will be generated, subject to `-z` limitations. Table names are separated by whitespace.

-td

specifies a statement delimiter for SQL statements generated by `dblook`. If a statement delimiter option is not specified, the default is the semicolon (';'). At the end of each DDL statement, the delimiter is printed, followed by a new line.

-o <filename>

specifies the file where the generated DDL is written. If this file is not specified, it defaults to the console (i.e. standard System.out).

-append

prevents overwriting the DDL output ('-o' parameter, if specified) and "dblook.log" files. If this parameter is specified, and execution of the `dblook` command leads to the creation of files with names identical to existing files in the current directory, `dblook` will append to the existing files. If this parameter is not set, the existing files will be overridden.

-verbose

specifies that all errors and warnings (both SQL and internal to `dblook`) should be echoed to the console (via `System.err`), in addition to being printed to the "dblook.log" file. If this parameter is not set, the errors and warnings only go to the "dblook.log" file.

-noview

specifies that CREATE VIEW statements should not be generated.

Generating the DDL for a database

The `dblook` utility generates all of the following objects when generating the DDL for a database:

- Checks
- Functions
- Indexes
- Jar files
- Keys (primary, foreign, and unique)
- Schemas
- Stored procedures
- Triggers
- Tables
- Views

Note: When `dblook` runs against a database that has jar files installed, it will create a new directory, called `DERBYJARS`, within the current directory, and that is where it will keep copies of all of the jars it encounters. In order to run the generated DDL as a script, this `DERBYJARS` directory must either 1) exist within the directory in which it was created, or 2) be moved manually to another directory, in which case the path in the generated DDL file must be manually changed to reflect to the new location.

The `dblook` utility ignores any objects that have system schemas (for example, `SYS`, `SYSIBM`), since DDL is not able to directly create nor modify system objects.

dblook examples

The following examples demonstrate how the various `dblook` utility options might be specified from a command line. Each example is preceded by a brief description of how the `dblook` utility behaves given the specified options.

Dump the DDL for everything in database 'sample' (in the current directory) to the console.

```
java org.apache.derby.tools.dblook -d jdbc:derby:sample
```

Dump the DDL for everything in database 'sample' (in the current directory) to the console, and include error/warning messages (with the latter being printed via `System.err`)

```
java org.apache.derby.tools.dblook
-d jdbc:derby:sample -verbose
```

Dump the DDL for everything in database 'sample' (in the current directory)
to a file called "myDB_DDL.sql" (in the current directory).
java org.apache.derby.tools.dblook -d jdbc:derby:sample -o myDB_DDL.sql

Dump the DDL for everything in database 'sample' (in a specified directory)
to a file called "newDB.sql" (in a specified directory).
java org.apache.derby.tools.dblook -d
'jdbc:derby:c:\private\stuff\sample'
-o 'C:\temp\newDB.sql'

Dump the DDL for all objects in database 'sample' with schema 'SAMP' to the console.
java org.apache.derby.tools.dblook -d jdbc:derby:sample -z samp

Dump the DDL for all objects in remote database 'sample' on 'localhost:1527' with schema 'SAMP' to the console.
java org.apache.derby.tools.dblook -d
'jdbc:derby://localhost:1527/"C:\temp\sample";
user=someusr;password=somepwd' -z samp

Dump the DDL for all objects with schema 'SAMP' in database 'sample' that are related to the table name 'My Table' to the console.
java org.apache.derby.tools.dblook -d jdbc:derby:sample -z samp -t "My Table"

Dump the DDL for all objects in database 'sample' that are related to either of the table names 'STAFF' or 'My Table' to the console.
java org.apache.derby.tools.dblook -d jdbc:derby:sample -t "My Table" staff

Dump the DDL for all objects in database 'sample' to myDB_DDL.sql, with no statement delimiter (i.e. leave off the default semi-colon), and append to the output files if they are already there.
java org.apache.derby.tools.dblook -d jdbc:derby:sample
-o myDB_DDL.sql -td '' -append

Dump the DDL for all objects EXCEPT views in database 'sample' to the console.
java org.apache.derby.tools.dblook -d jdbc:derby:sample -noview

Status messages are printed to the output (either a -o filename, if specified, or the console) as SQL script comments. These status messages serve as headers to show which types of database objects are being (or have been) processed by the dblook utility.

Trademarks

The following terms are trademarks or registered trademarks of other companies and have been used in at least one of the documents in the Apache Derby documentation library:

Cloudscape, DB2, DB2 Universal Database, DRDA, and IBM are trademarks of International Business Machines Corporation in the United States, other countries, or both.

Microsoft, Windows, Windows NT, and the Windows logo are trademarks of Microsoft Corporation in the United States, other countries, or both.

Java and all Java-based trademarks are trademarks of Sun Microsystems, Inc. in the United States, other countries, or both.

UNIX is a registered trademark of The Open Group in the United States and other countries.

Other company, product, or service names may be trademarks or service marks of others.